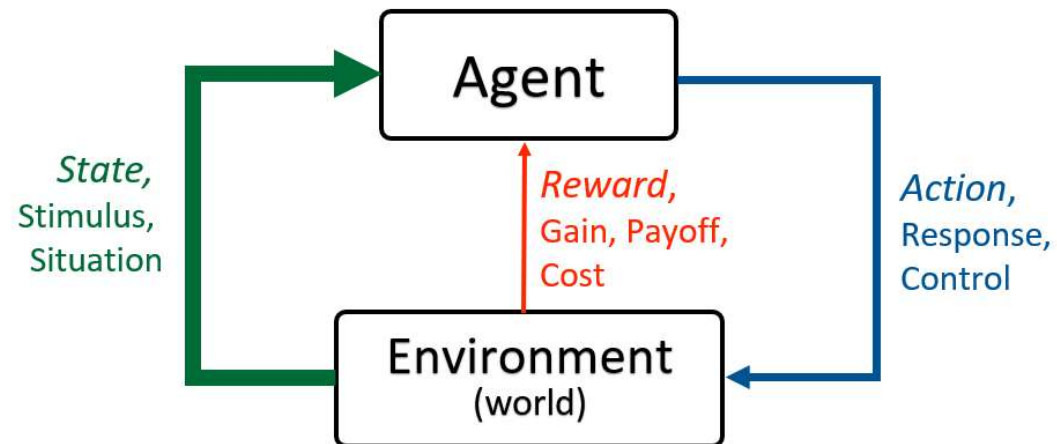




Reinforcement Learning 1

Anders Lyhne Christensen




What is the reward?



Challenges

- An agent that typically has **no prior knowledge** must learn how to solve a task through interaction with an environment.
- The agent may receive a reward after each action, or only from time to time – perhaps only very infrequently.



Example: Learning to play football from scratch is challenging: when should an agent be rewarded?

Sparse and deceptive rewards can complicate learning.

Multiarmed/ k -armed Bandit Problem

Multiarmed/ k -armed Bandit Problem

Suppose there are k bandits.

Action a_i is like pulling a slot machine arm with random payoff function $R(a_i)$



$R(a_1)$



$R(a_2)$

...



...

$R(a_k)$

Definitions

$$q_*(a) \doteq \mathbb{E}[R_t \mid A_t = a]$$

Expected reward

Arbitrary action

Expected value

Reward at time t

Action selected at time t

Assume that there are $k=4$ bandits

- Action: pull arm 1 — Reward is always 8

- value of arm 1 is $q_*(1) =$


- Action: pull arm 2 — 88% chance of 0 and 12% chance of 100!

- value of arm 2 is $q_*(2) = .88 \times 0 + .12 \times 100 =$


- Action: pull arm 3 — Randomly between -10 and 35, equiprobable



- Action: pull arm 4 — a third 0, a third 20, a third from $\{8,9,\dots, 18\}$



We don't know $q_*(a)$, but we want to estimate it. We call our estimate



The k -armed Bandit Problem

- On each of a sequence of *time steps*, $t=1, 2, 3, \dots$, you choose an action A_t from k possibilities, and receive a real-valued *reward* R_t
- The reward depends only on the action taken; it is identically, independently distributed (i.i.d.):

$$q_*(a) \doteq \mathbb{E}[R_t | A_t = a], \quad \forall a \in \{1, \dots, k\} \quad \text{true values}$$

- These true values are *unknown*. The distribution is unknown.
- Nevertheless, you must maximize your total reward
- You must both try actions to learn their values (*explore*), and prefer those that appear best (*exploit*)

Bounded Reward Assumption

A common assumption we will make is that rewards are in a bounded interval $[-R_{max}, R_{max}]$.

I.e., for each i , $\Pr(R(a_i) \in [-R_{max}, R_{max}]) = 1$.

Note that results are available for other types of assumptions, e.g. Gaussian distributions of rewards.

Multi-Armed Bandits: Application examples

- **Clinical Trials**

- Arms: possible treatments
- Arm pulls: application of treatment to individual
- Rewards: outcome of treatment
- Objective: maximize cumulative reward = maximize benefit to trial population (or find best treatment quickly)

- **Online Advertising**

- Arms: different ads/ad-types for a webpage
- Arm pulls: displaying an ad upon a page access
- Rewards: click through
- Objective: maximize cumulative reward = maximize clicks (or find best add quickly)

Assuming that you have k bandits, how would you maximize the rewards?



Remember: bandits pay out a reward chosen from probability distributions. Each bandit has its own probability distribution, and you have no idea what they are.

UniformBandit Algorithm

1. Pull each arm a fixed number (w) times
2. Return arm with best average reward



$r_{11} \ r_{12} \ \dots \ r_{1w}$

$r_{21} \ r_{22} \ \dots \ r_{2w}$

$r_{k1} \ r_{k2} \ \dots \ r_{kw}$

Imagine that we have to
maximize the reward over a total
of 1000 arm pulls.

What should we do?

The Exploration/Exploitation Dilemma

- Suppose you form estimates

$$Q_t(a) \approx q_*(a), \quad \forall a \quad \text{action-value estimates}$$

- Define the *greedy action* at time t as

$$A_t^* \doteq \arg \max_a Q_t(a)$$

- If $A_t = A_t^*$ then you are *exploiting*
If $A_t \neq A_t^*$ then you are *exploring*
- You can't do both at the same time.

Action-Value Methods

- Methods that learn **action-value estimates** and nothing else
- For example, estimate action values as *sample averages*:

$$Q_t(a) \doteq \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot \mathbb{1}_{A_i=a}}{\sum_{i=1}^{t-1} \mathbb{1}_{A_i=a}}$$

- The sample-average estimates converge to the true values *if* the action is taken an infinite number of times

$$\lim_{N_t(a) \rightarrow \infty} Q_t(a) = q_*(a)$$

The number of times action a
has been taken by time t



ϵ -Greedy Action Selection

- In **greedy action selection**, you always exploit
- In **ϵ -greedy**, you are usually greedy, but with probability ϵ you instead pick an action at random (possibly the greedy action again)
- This is perhaps the simplest way to balance exploration and exploitation

Assume that an action has been taken $n - 1$ times already. Our estimated reward of taking the action an n 'th time is then:

$$Q_n = \frac{R_1 + R_2 + \dots + R_{n-1}}{n - 1}$$

Not computationally
efficient



$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} \left(R_n + \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} \left(R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} (R_n + (n-1)Q_n) \\ &= \frac{1}{n} (R_n + nQ_n - Q_n) \\ &= Q_n + \frac{1}{n} [R_n - Q_n], \end{aligned}$$

A simple bandit algorithm

Initialize, for $a = 1$ to k :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Repeat forever:

$$A \leftarrow \begin{cases} \arg \max_a Q(a) & \text{with probability } 1 - \varepsilon \\ \text{a random action} & \text{with probability } \varepsilon \end{cases} \quad \text{(breaking ties randomly)}$$

$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$

$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$

Averaging as the learning rule

This is a standard form for learning/update rules:

$$NewEstimate \leftarrow OldEstimate + StepSize \underbrace{[Target - OldEstimate]}_{\text{Error}}$$

Potentially noisy

Learning rate α


The diagram shows the equation $NewEstimate \leftarrow OldEstimate + StepSize [Target - OldEstimate]$. A handwritten orange arrow points from the text 'Potentially noisy' to the 'Target' term. Another handwritten orange arrow points from the text 'Learning rate α ' to the 'StepSize' term. A third handwritten orange bracket is drawn under the expression '[Target - OldEstimate]', with the word 'Error' written below it.

Note, that the “*Target*” is presumed to indicate a desirable direction in which to move (may be noisy). The “*Target*” here is the n ’th reward.



How well does ε -greedy action selection work?

Let's test for different values of ε .

The 10-armed Testbed



ϵ -Greedy Methods on the 10-Armed Testbed



Tracking a Non-stationary Problem

- Suppose the true action values **change slowly over time** - then we say that the problem is *non-stationary*
- In this case, sample averages are not a good idea (Why?)
- Better is an “*exponential, recency-weighted average*”:

$$Q_{n+1} \doteq Q_n + \alpha [R_n - Q_n]$$


Where α , the step-size or learning rate, is constant.

- What is the effect if setting $\alpha = 1$?
- What is the effect if setting $\alpha = 0$?
- Which value should we assign to α ?

We will still focus on stationary problems, but remove the initial bias over time by setting α to a constant value.

Optimistic Initial Values


- All methods so far depend on $Q_1(a)$, i.e., they are biased. So far, we have used $Q_1(a) = 0$.
- Suppose we initialize the action values *optimistically* ($Q_1(a) = 5$), on the 10-armed testbed (with $\alpha = 0.1$)



Upper Confidence Bound (UCB) action selection

- A clever way of reducing exploration over time
- Estimate an upper bound on the true action values
- Select the action with the largest (estimated) upper bound

$$A_t \doteq \operatorname{argmax}_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$




Gradient Bandit Algorithm

- We don't try to estimate or rewards, we just keep track of our preference (or probability) for choosing an action:


$$\Pr\{A_t = a\} \doteq \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} \doteq \pi_t(a)$$

- We learn preferences according to:

$$\begin{aligned} H_{t+1}(A_t) &\doteq H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)), & \text{and} \\ H_{t+1}(a) &\doteq H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), & \text{for all } a \neq A_t, \end{aligned}$$



Results



Summary 1/2

- *Reinforcement learning*: an agent learns through interactions with an environment.
- *k*-armed bandit problem: The true distribution of rewards is not known a priori, but must be learned.



$$q_*(a) \doteq \mathbb{E}[R_t \mid A_t = a]$$

$$Q_t(a) \approx q_*(a), \quad \forall a$$

Summary 2/2

We discussed:

- UniformBandit
- Exploration vs. exploitation tradeoff
- ε -Greedy
- How to reduce the bias of initial estimates $Q_1(a)$, namely a fixed step-size α .
- Optimistic initialization ($Q_1(a) = 5$)
- Upper confidence bound: $A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$
- Gradient bandit algorithm: preferences are used instead of estimates of q_*